

Inference throughput: master + worker pool

