

Inference throughput: master + worker pool

LLM queries / s (log scale)

